

Blake Bullwinkel

CONTACT INFORMATION

✉ blakebullwinkel@gmail.com
🌐 blakebullwinkel.com

in linkedin.com/in/blakebullwinkel
🐙 github.com/blakebullwinkel

EDUCATION

Harvard University, Cambridge, MA May 2022
M.S. in Data Science. GPA: 3.95/4
Advisor: Pavlos Protopapas
Thesis: *Generative Adversarial Network Methods for Solving Differential Equations*

Williams College, Williamstown, MA June 2020
B.A. in Mathematics, Chinese. GPA: 3.83/4 (*cum laude*)

University of Oxford, Oxford, UK June 2019
Attended as part of the selective, year-long Williams-Exeter Program at Oxford.

PUBLICATIONS

R Pellegrin*, **B Bullwinkel***, M Mattheakis, P Protopapas. *Transfer Learning with Physics-Informed Neural Networks for Efficient Simulation of Branched Flows*. NeurIPS Workshop on Machine Learning and the Physical Sciences, 2022.

B Bullwinkel*, D Randle*, P Protopapas, D Sondak. *DEQGAN: Learning the Loss Function for PINNs with Generative Adversarial Networks*. ICML Workshop on AI for Science (AI4Science), 2022.

B Bullwinkel, K Grabarz, L Ke, Sc Gong, C Tanner, J Allen. *Evaluating the Fairness Impact of Differentially Private Synthetic Data*. ICML Workshop on Theory and Practice of Differential Privacy (TPDP), 2022.

RESEARCH EXPERIENCE

AI Safety and Alignment, Harvard University Sept 2023–Present
Fall Research Course. Advisors: Weiwei Pan, Finale Doshi-Velez, Claude Bruderlein

- Leading a team of graduate students to characterize and quantify emerging failure modes of LLMs (including hallucinations, value misalignment, and adversarial attacks) in the context of an LLM-based simulator for humanitarian negotiation training.

Multi-Modal Adversarial Attacks, Harvard University Sept 2023–Present
Fall Research Course. Advisors: Siddarth Swaroop, Weiwei Pan, Finale Doshi-Velez

- Characterizing the properties of adversarial attacks on LLMs and Vision Language Models (VLMs) that exploit soft prompting.

Physics-Informed Neural Networks, Harvard University Feb 2021–May 2022
Master's Thesis. Advisors: Pavlos Protopapas, David Sondak

- Researched methods to improve the training stability of DEQGAN, a generative adversarial network for obtaining accurate solutions to a wide range of ordinary and partial differential equations.
- Developed multi-head architectures and transfer learning algorithms to more efficiently simulate branched flows, a universal wave phenomenon.
- Maintained research code in a user-friendly [PyTorch package](#).

Interpretable Machine Learning, Harvard University Feb 2022–May 2022
Spring Research Course. Advisors: Weiwei Pan, Yaniv Yacoby

- Investigated how non-identifiability in additive models can cause misleading model interpretations in the healthcare domain.
- Characterized a particular form of non-identifiability that arises when generalized additive models are trained on data with interaction effects.

Differential Privacy and Fairness, Microsoft Sept 2021–Dec 2021
IACS Capstone Project. Advisor: Joshua Allen

- Led a collaboration among graduate students and Microsoft researchers to understand the fairness impact of training machine learning models on differentially private synthetic data.
- Proposed a simple pre-processing technique to synthesize data that promote more fair outcomes without degrading accuracy.

Feb 2020

- Applied compartmental models to early COVID-19 data published by the Chinese National Health Commission to estimate key disease parameters and simulate an outbreak on a college campus with a quarantine policy.

Aug 2022–Present

- Deployed an LLM-powered app that answers questions about internal documentation using retrieval augmented generation (Azure OpenAI Service, Azure Cognitive Search, LangChain).
- Correlated performance bugs to identify common root causes by creating a vector database of text embeddings.
- Built a pipeline to detect kernel-mode memory leaks across the Azure fleet and worked with OS engineers to mitigate the highest impact leaks, saving around 2TB of memory per day (project received Quality Stars award for FY23 Q3).
- Trained ML models that help deployment teams assess the risk of Azure Host OS updates.

Summer 2021

- Developed a Python package for anomaly detection of water usage time series data.
- Trained models (ARIMA, Facebook Prophet, LSTM) to forecast water usage efficiency, lowering the mean validation RMSE by 43% in comparison to a moving average baseline.
- Developed an automated data pipeline with actionable insights in Power BI, adopted by beverage plants nationwide.

June 2020–Jan 2022

- Led the development of an iOS **mobile app** that provides carbon footprint estimates for grocery products.
- Built Google Firebase backend with 150,000+ products scraped from supermarket websites.
- Accepted into the Harvard i-lab Venture Program for three consecutive semesters.

Summer 2019

- Performed time series analysis to identify and explain transaction delays in solar electricity startup's software platform.

Feb 2022–May 2022

- CS 109b: Advanced Topics in Data Science
- Prepared teaching materials and held office hours for students studying non-linear statistical methods and deep learning models, including CNNs, RNNs, LSTMs, autoencoders, GANs, and transformers.

2017–2020

- CHIN 201: Intermediate Chinese I (Fall 2017)
- CHIN 202: Intermediate Chinese II (Spring 2018)
- CHIN 301: Upper-Intermediate Chinese I (Fall 2019)
- CHIN 302: Upper-Intermediate Chinese II (Spring 2020)
- In 1:1 sessions, met weekly with students for casual discussions to practice spoken language, review vocabulary, and learn grammar structures.

August 2023–Present

- Leading lab sessions for AP Computer Science Principles at Global Impact Academy STEM Magnet High School.

Jan 2022

- Worked alongside professors to run workshop focused on teaching fundamental data science skills, including Python programming, probability theory, linear algebra, and statistics.

HONORS & AWARDS	Certificate of Distinction in Teaching , Harvard University Awarded based on student ratings (mean 4.67/5) for teaching of CS 109b.	2022
	IACS Student Scholarship , Harvard University Awarded to support data science thesis research at IACS (\$20,000 award).	2021
	Goldberg Prize in Mathematics , Williams College Awarded to the graduating senior who delivers the best mathematics colloquium.	2020
	Linen Senior Prize in Chinese , Williams College Awarded to the top graduating Chinese major.	2020
	Putnam Competition , MAA Scored 18, ranking fifth at Williams and in the top 15% of all participants.	2019
	Carolyn Altes Scholarship , AWCA Awarded on the basis of academics and potential to contribute to society.	2019
	Linen Grant , Williams College Awarded on the basis of academics to support summer study in China.	2017
	Davis UWC Scholar , Davis United World College Scholars Program Awarded to recognize commitment to building cross-cultural understanding.	2016
	Class of '16 Student Speaker , UWCSEA East Elected by peers to deliver the Class of '16 graduation student address.	2016
SKILLS & INTERESTS	Programming: Python (NumPy, pandas, sklearn, TensorFlow, PyTorch), R, SQL, KQL, HTML/CSS, JavaScript	
	Tools/Platforms: Conda, Jupyter, Git, Docker, Kubernetes, Azure, AWS	
	Language: Working proficiency in written and spoken Chinese (Mandarin)	
	Interests: Running, rowing, piano, writing (Medium blog), Rubik's cube solving (WCA profile)	
REFERENCES	Dr. Pavlos Protopapas Harvard University Email: pavlos@seas.harvard.edu	
	Dr. Weiwei Pan Harvard University Email: weiweipan@g.harvard.edu	
	Dr. Mihai Stoiciu Williams College Email: mstoiciu@williams.edu	
	Dr. Julie Blackwood Williams College Email: jcb5@williams.edu	